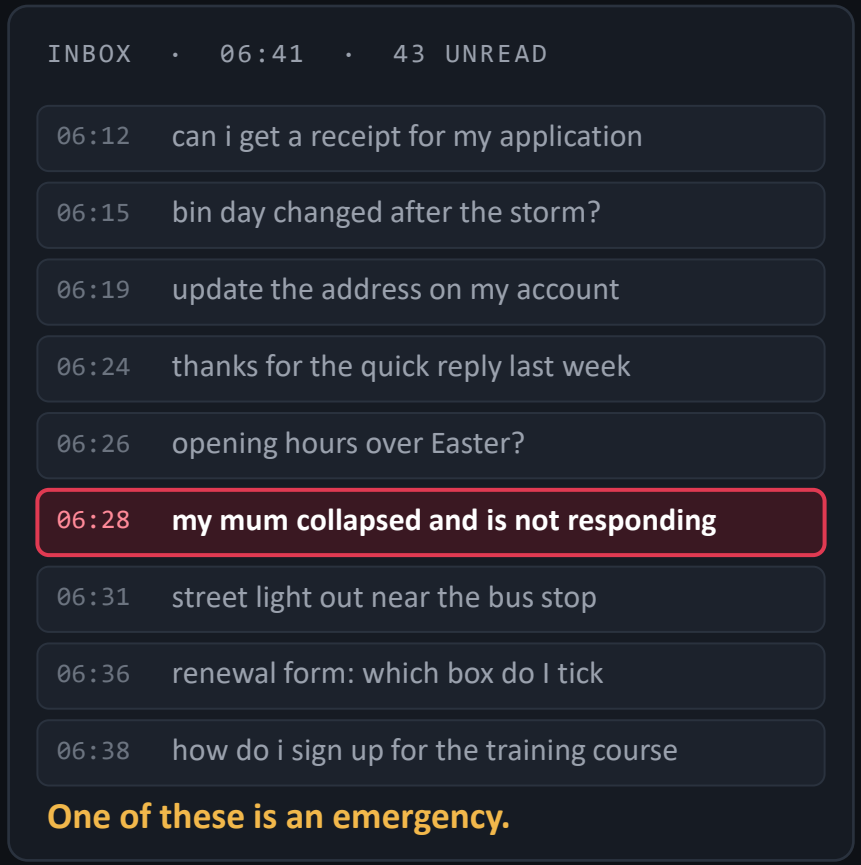
CODE FAIR 2026 · AI CHALLENGE 5

Every urgent message reaches a person.

NT Correspondence Triage. It sorts, routes and explains every message, and it never replies by itself.

TEAM AIC023

NT CORRESPONDENCE TRIAGE



THE NORTHERN TERRITORY, WET SEASON

A short message, from a phone with one bar of signal.

From October to April the monsoon can cut roads and take out phone towers. In March 2024 it did both at Borroloola, and mobile coverage across the area went down.

Sources: ABS, RMIT / ADM+S, AAP

58.5%

of First Nations people speak an Indigenous language at home

ABS, 2021 Census

43%

of First Nations communities and homelands have no mobile service

RMIT / ADM+S, 2023

3 in 4

First Nations people in remote communities are digitally excluded

RMIT / ADM+S, 2025

4.4×

the national rate of assault victims (4,057 vs 917 per 100,000)

ABS, 2025

Triage. Escalate. Never reply.



Message arrives

SMS, web form, email or phone notes

Three readers

AI model, backup ML model, safety rules

Priority only rises

Floors and access can raise it, never lower it

A named person

Right team, Duty Officer copied on critical

A clock runs

15, 60 or 240 minutes, then it escalates

The system drafts. A person approves. It never sends anything to the public.

Three readers. None can lower a priority.

The AI model

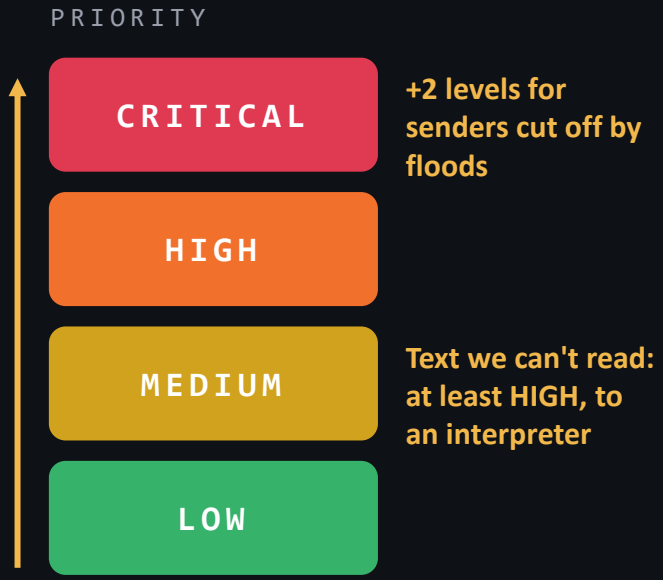
Reads the meaning. Typos and slang don't lower urgency. Treats every message as data, not orders.

The backup ML model

Trained on our data. Runs with no network, and its disagreement is a warning.

The safety rules

61 fixed patterns, from "hurt me" to "not breathing". They ignore instructions in the text.



WE TESTED IT, FOUND A GAP, FIXED IT

The second-opinion floor

BENCHMARK MESSAGE M027 · TRULY HIGH

“fence and half the carport gone in the wind last night, water starting to pool near the house”

AI MODEL SAID
MEDIUM

≠

BACKUP MODEL SAID
HIGH

The system had been throwing that disagreement away, so a high message was missed.

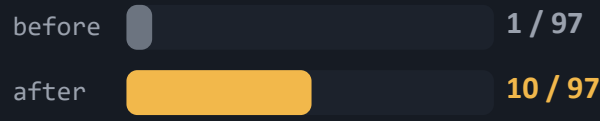
RULE backup says high or critical, AI lower: at least HIGH until a person reads it

WHAT IT CHANGED · 152 MESSAGES

Urgent messages caught



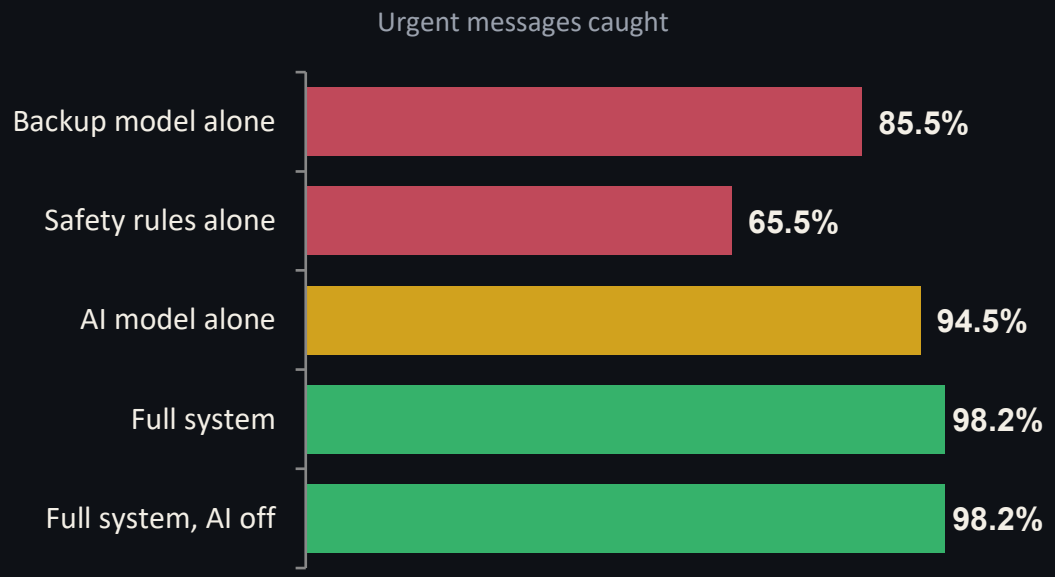
Routine messages escalated



The price: 9 more routine messages to read. We think that trade is right.

Does it work?

152 HAND-WRITTEN MESSAGES · REAL AI VERDICTS · 55 TRULY URGENT



54 / 55 urgent messages caught

23 / 23 critical messages on the 15-minute clock

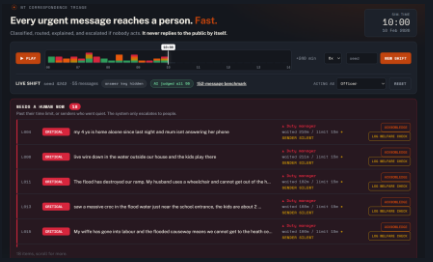
10 / 97 routine messages escalated: the price of safety

Synthetic data that we wrote. Real-world scores will be lower.

Fairness, language and a late bug >

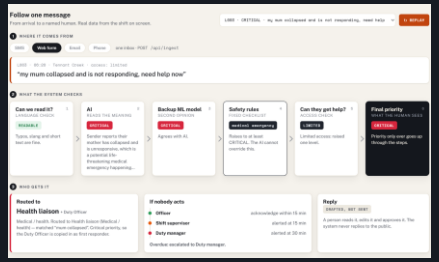
What is inside the console

REAL SCREENSHOTS FROM THE RUNNING SYSTEM



Shift timeline

Drag the clock. The dark panel lists what is overdue or silent.



Follow one message

Six checks, then the named team and the escalation ladder.



NT map

Messages flow to the triage centre. Pins pulse when overdue.



Live scores

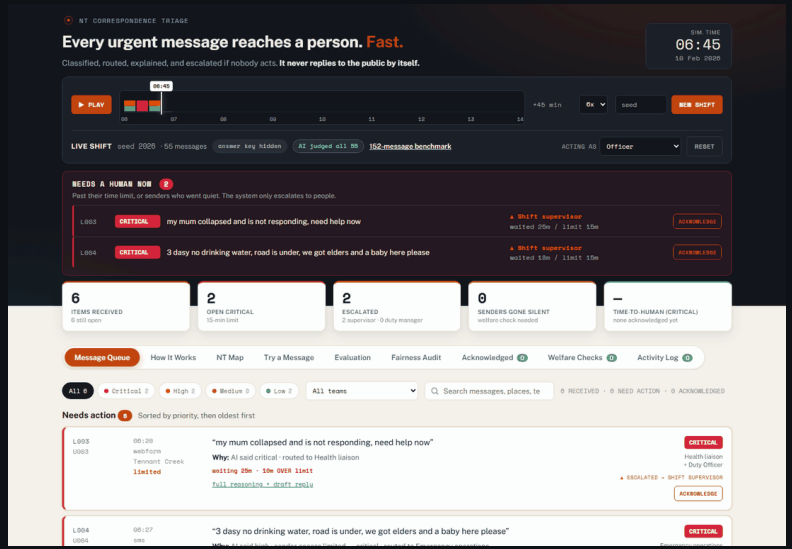
Computed against a hidden answer key, with the AI switched off too.

ALSO: THREE ROLES · WELFARE CHECKS · AUDIT LOG · TRY ANY MESSAGE

Live demo

CLICK A STEP TO OPEN IT

1. The live queue, and who needs a human now
2. Follow one message through the checks
3. Try to break it: injection, unreadable text
4. Nobody acts, so it escalates
5. The scores, even with the AI off



OPEN THE CONSOLE ▶

127.0.0.1:5050 · seed 2026 · no AI credit used

What is real. What is next. What we won't do.

Real today

The full pipeline, escalation clocks, silent-sender checks, an inbound endpoint, an audit trail and 97 tests. AI verdicts are real.

Next

A shadow pilot for one wet season. Real SMS and email feeds. Language phrases reviewed by speakers, under community control.

Never

No auto-replies. No lowering a priority. No pretending to understand a language we can't read.

OUR WARNINGS: SYNTHETIC DATA · 30% ON THE CRITICAL CLOCK · ONE MODEL, ENGLISH ONLY

Questions?

Help people. Never replace them.

Likely questions >